

Week 3 — Task 1: Trained Machine Learning Models

Dataset: RIT Opportunity Wise Dataset (8,558 rows, 29 columns)

Target Variable: `Is_Dropped` — 1 if the student dropped out or withdrew, 0 otherwise

Models Used: Logistic Regression (baseline) + Random Forest Classifier (main model)

Step 1: Import Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.preprocessing import LabelEncoder
print("libraries imported successfully.")
```

Step 2: Load Dataset

```
df = pd.read_csv(r"T:\Jupyter\RIT_Dataset_Week2_Cleaned.csv")
print("Shape:", df.shape)
df.head(3)
```

Shape:
(8558,
29)

	Learner	SignUp	DateTime
Opportunity Id \			
0	2023-06-14		
12:30:35			
00000000-0GN2-			
A0AY-7XK8-C5FZP			
1	2023-05-01		
05:29:16			
00000000-0GN2-			
A0AY-7XK8-C5FZP			
2	2023-04-09		
20:35:08			
00000000-0GN2-			
A0AY-7XK8-C5FZP			

Opportunity Name
Opportunity
Category

\

0 Career
Essentials:
Getting Started
with Your P...
Cours

1 Career
Essentials:
Getting Started
with Your P...
Cours

2 Career
Essentials:
Getting Started
with Your P...
Cours

Opportunity	End Date	First Name	Date of Birth	Gender	Country
-------------	----------	------------	---------------	--------	---------

\

0	2024-06-29 18:52:39	Faria	2001-01-12	Female	Pakistan
---	---------------------	-------	------------	--------	----------

1	2024-06-29 18:52:39	Poojitha	2000-08-16	Female	India
---	---------------------	----------	------------	--------	-------

2	2024-06-29 18:52:39	Emmanuel	2002-01-27	Male	United States
---	---------------------	----------	------------	------	---------------

0	Institution Name	...	Signup Year	Signup Month	\
	Nwihs	...	2023.0	6.0	

1	SAINT LOUIS	...	2023.0	5.0	
---	-------------	-----	--------	-----	--

2	Illinois Institute of Technology	...	2023.0	4.0	
---	----------------------------------	-----	--------	-----	--

Cohort_Missing	Signup Cohort	Days to Apply	Apply to Start	Gap	\
0	False	2023-06	0.0	-223.0	
1	False	2023-05	0.0	-179.0	
2	False	2023-04	NaN	NaN	

```

    Opportunity Duration (Days) Engagement Stage Is_Completed 0
Is_Droppe 0
    604.0 Active 0 1
    0
    604.0 Active 0 2
    0
    604.0 Active 0
[3 rows x
29
columns]

```

Step 3: Define Target & Drop Leakage Columns

The target is `Is_Dropped`. Columns dropped here fall into four categories:

- **Data leakage** — derived from the target itself (`Status Description`, `Status Code`, `Engagement Stage`, `Is_Completed`)
- **Identifiers** — names, IDs, and institution details that carry no predictive signal
- **Raw date columns** — already replaced by engineered features (`Days to Apply`, `Apply to Start Gap`)
- **Low-signal flags** — `Cohort_Missing`, `Start_Date_Missing`, `Opportunity Duration (Days)`

```

drop_cols = [
    #
    target + leakage
    'Is_Dropped', 'Is_Completed', 'Engagement Stage',
    'Status Description', 'Status Code',
    # identifiers
    'First Name', 'Opportunity Id', 'Opportunity Name',
    'Institution Name', 'Current/Intended Major', # raw
    dates (replaced by engineered features)
]

```

```

    'Learner SignUp DateTime', 'Apply Date',
    'Opportunity Start Date', 'Opportunity End Date', 'Date of Birth',
    'Entry created at',
    # low-signal / flag columns
    'Signup Cohort', 'Cohort_Missing', 'Start_Date_Missing',
    'Opportunity Duration (Days)',
]

y = df['Is_Dropped']
X = df.drop(columns=drop_cols)

print("Target distribution:")
print(y.value_counts())
print("\nFeatures used for training:")
print(X.columns.tolist())

Target distribution:
Is_Dropped
0          7855
1           703
Name: count, dtype: int64
Features used for training:
['Opportunity Category', 'Gender', 'Country', 'Age', 'Age Group',
'Signup Year', 'Signup Month', 'Days to Apply', 'Apply to Start Gap']

```

Step 4: Label Encoding + Null Handling

ML models require numeric inputs. All categorical text columns (Gender, Country, Opportunity Category, Age Group) are converted using `LabelEncoder`, which assigns each unique string a unique integer. Boolean columns are cast to 0/1. Remaining nulls — mainly in `Days to Apply` and `Apply to Start Gap` due to missing date records — are filled with the column median, which preserves the distribution without introducing outliers.

```

x = X.copy()

# encode categorical text columns for col in
x.select_dtypes(include='object').columns:      x[col] =
LabelEncoder().fit_transform(x[col].astype(str))

# encode boolean columns for col in
x.select_dtypes(include='bool').columns:      x[col]
= x[col].astype(int)

# fill remaining nulls with median
x = x.fillna(x.median())

print("Encoded feature matrix shape:", x.shape)
print(x.dtypes)

```

```
Encoded feature matrix shape: (8558, 9)
Opportunity Category      int64
Gender                    int64
Country                   int64
Age                       int64
Age Group                  int64
Signup Year                float64
Signup Month               float64
Days to Apply              float64
Apply to Start Gap         float64
dtype: object
```

Step 5: Train-Test Split

Data is split 80/20 — 80% for training, 20% held out for testing. `random_state=42` ensures the split is reproducible and identical for both models, making the accuracy comparison fair.

```
x_train, x_test, y_train, y_test = train_test_split(
    x, y, test_size=0.2, random_state=42
)

print("Training set:", x_train.shape)
print("Test set      :", x_test.shape)
print("\nClass balance in training set:")
print(y_train.value_counts())
Training set: (6846, 9)
Test set      : (1712, 9)
Class balance in training set:
Is_Dropped
0         6297
1          549
Name: count, dtype: int64
```

Model 1: Logistic Regression (Baseline)

Algorithm: Logistic Regression draws a linear decision boundary between the two classes based on a weighted combination of input features. It outputs a dropout probability via the sigmoid function and classifies a student as "Dropped" if that probability exceeds 0.5.

Key Parameters:

- `max_iter=1000` — allows the solver to fully converge given the imbalanced data
- `class_weight='balanced'` — automatically adjusts class weights inversely proportional to frequency; without this, the model ignores the minority Dropped class entirely (only ~8% of data)
- `random_state=42` — reproducibility

Why use it: Logistic Regression is an interpretable baseline. It tells us the minimum acceptable performance that a more complex model must beat to justify its extra complexity.

```
lr = LogisticRegression(max_iter=1000, class_weight='balanced',
random_state=42)
lr.fit(x_train, y_train)

lr_pred = lr.predict(x_test)
lr_accuracy = accuracy_score(y_test, lr_pred)

print("Logistic Regression Accuracy:", round(lr_accuracy * 100, 2),
"%")
print("\nClassification Report:")
print(classification_report(y_test, lr_pred, target_names=['Not
Dropped', 'Dropped']))
```

Logistic Regression Accuracy: 63.49 %
Classification Report:

	precision	recall	f1-score	support
Not Dropped	0.96	0.62	0.76	1558
Dropped	0.17	0.77	0.27	154
accuracy			0.63	1712
macro avg	0.57	0.69	0.52	1712
weighted avg	0.89	0.63	0.71	1712

Model 2: Random Forest Classifier (Main Model)

Algorithm: Random Forest is an ensemble of 100 decision trees. Each tree is trained on a random bootstrap sample of the data and evaluates only a random subset of features at each split. At prediction time, all 100 trees vote and the majority class wins. This makes it robust to overfitting and capable of capturing non-linear interactions between features.

Key Parameters:

- `n_estimators=100` — 100 trees; enough for stable and accurate predictions
- `class_weight='balanced'` — same minority-class correction as Logistic Regression for fair comparison
- `random_state=42` — reproducibility

Why it is expected to outperform Logistic Regression: Dropout likelihood is influenced by combinations of factors — e.g., young students applying to internships with long gaps before start — that a single linear boundary cannot capture. Random Forest models these interactions naturally.

```
rf = RandomForestClassifier(n_estimators=100, class_weight='balanced',
random_state=42)
```

```

rf.fit(x_train, y_train)

rf_pred = rf.predict(x_test)
rf_accuracy = accuracy_score(y_test, rf_pred)

print("Random Forest Accuracy:", round(rf_accuracy * 100, 2), "%")
print("\nClassification Report:")
print(classification_report(y_test, rf_pred, target_names=['Not
Dropped', 'Dropped']))

```

Random Forest Accuracy: 87.73 %
Classification Report:

	precision	recall	f1-score	support
Not Dropped	0.95	0.91	0.93	1558
Dropped	0.37	0.50	0.42	154
accuracy			0.88	1712
macro avg	0.66	0.71	0.68	1712
weighted avg	0.90	0.88	0.89	1712

Model Comparison


```

comparison = pd.DataFrame({
    'Model': ['Logistic Regression', 'Random Forest'],
    'Accuracy (%)': [
        round(lr_accuracy * 100, 2),
        round(rf_accuracy * 100, 2)
    ]
})
print(comparison)

fig, ax = plt.subplots(figsize=(6, 4))
bars = ax.bar(comparison['Model'], comparison['Accuracy (%)'],
color=['#2E75B6', '#C00000'], width=0.4) ax.set_title('Model Accuracy
Comparison', fontsize=13,
fontweight='bold')
ax.set_ylabel('Accuracy (%)') ax.set_ylim(0, 100) for bar in bars:
ax.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
f'{bar.get_height():.2f}%', ha='center', fontweight='bold')
plt.tight_layout()
plt.show()
print("Random Forest selected as the primary model for Tasks 2 and
3.")

```

	Model	Accuracy (%)
0	Logistic Regression	63.49
1	Random Forest	87.73



Random Forest selected as the primary model for Tasks 2 and 3.

TASK 2: PREDICTED PROBABILITIES

STEP 1: Generate Predicted Probabilities on Test Set

```
rf_proba = rf.predict_proba(x_test)
prob_not_dropped = rf_proba[:, 0]  prob_dropped
= rf_proba[:, 1]

print("-" * 70)
print("TASK 2: PREDICTED PROBABILITIES") print("-"
" * 70)
print(f"\nProbability Statistics (Test Set - {len(x_test)}
students):")
print(f"  Participation Strength (1 - Dropout Risk):")
print(f"    Mean:    {prob_not_dropped.mean():.4f}")
print(f"    Min:     {prob_not_dropped.min():.4f}")
print(f"    Max:     {prob_not_dropped.max():.4f}")
print(f"    Median:  {np.median(prob_not_dropped):.4f}")

print(f"\n  Dropout Risk Probability:") print(f"
Mean:    {prob_dropped.mean():.4f}") print(f"
Min:     {prob_dropped.min():.4f}") print(f"    Max:
{prob_dropped.max():.4f}") print(f"    Median:
{np.median(prob_dropped):.4f}")

-----

TASK 2: PREDICTED PROBABILITIES

-----

Probability Statistics (Test Set - 1712 students):
  Participation Strength (1 - Dropout Risk):
    Mean:    0.8706
    Min:     0.0000
    Max:     1.0000
    Median:  1.0000
  Dropout Risk Probability:
    Mean:    0.1294
    Min:     0.0000
    Max:     1.0000
    Median:  0.0000
```

STEP 2: Create Probabilities DataFrame with Opportunity Information

```
x_test_original = X.iloc[x_test.index].copy()
y_test_original = y.iloc[x_test.index].copy()

prob_df = pd.DataFrame({
    'Student_ID': range(len(x_test)),
    'Opportunity_Category': x_test_original['Opportunity
Category'].values,
    'Actual_Is_Dropped': y_test_original.values,
    'Predicted_Dropout_Risk': np.round(prob_dropped, 4),
    'Predicted_Participation_Strength': np.round(prob_not_dropped, 4),
})

prob_df['Participation_Category'] = pd.cut(
prob_df['Predicted_Participation_Strength'],
bins=[0, 0.3, 0.5, 0.7, 1.0],
labels=['Very Low (0-0.3)', 'Low (0.3-0.5)', 'High (0.5-0.7)',
'Very High (0.7-1.0)']
)

print("\n" + "-" * 70)
print("PARTICIPATION STRENGTH BY CATEGORY") print("-
" * 70)
print("\nTop 20 Students with Highest Participation Strength:")
```

```
print(prob_df.nlargest(20, 'Predicted_Participation_Strength')[      ['Opportunity_Strength',
'Predicted_Participation_Strength',
'Actual_Is_Dropped']
].to_string(index=False))
```

```
-----
PARTICIPATION STRENGTH BY CATEGORY
-----
```

Top 20 Students with Highest Participation Strength:

Opportunity_Category	Predicted_Participation_Strength	Actual_Is_Dropped
----------------------	----------------------------------	-------------------

Event	1.0	
-------	-----	--

Course	1.0	
--------	-----	--

Internship	1.0	
------------	-----	--

Internship	1.0	
------------	-----	--

Internship	1.0	
------------	-----	--

Internship	1.0	
------------	-----	--

Internship	1.0	
------------	-----	--

Course	1.0	
--------	-----	--

Competition	1.0	
-------------	-----	--

Internship	1.0	
------------	-----	--

Competition	1.0	
-------------	-----	--

Competition	1.0	
-------------	-----	--

Event	1.0	
-------	-----	--

Course	1.0	
--------	-----	--

Internship	1.0	
------------	-----	--

Internship	1.0	
------------	-----	--

Internship	1.0	0
Internship	1.0	0
Competition	1.0	1
Competition	1.0	0

STEP 3: Aggregate by Opportunity Category

```
opportunity_stats = prob_dfs.groupby('Opportunity_Category').agg({
    'Predicted_Participation_Strength': ['mean', 'std', 'min', 'max'],
    'Predicted_Dropout_Risk': 'mean',
    'Student_ID': 'count'
}).round(4)

opportunity_stats.columns = ['_'.join(col).strip() for col in
opportunity_stats.columns.values]
opportunity_stats.rename(columns={
    'Predicted_Participation_Strength_mean':
    'Avg_Participation_Strength',
    'Predicted_Participation_Strength_std': 'Std_Dev',
    'Predicted_Participation_Strength_min': 'Min_Strength',
    'Predicted_Participation_Strength_max': 'Max_Strength',
    'Predicted_Dropout_Risk_mean': 'Avg_Dropout_Risk',
    'Student_ID_count': 'Number_of_Students'
}, inplace=True)

opportunity_stats =
opportunity_stats.sort_values('Avg_Participation_Strength',
ascending=False)

print("\n" + "-" * 70)
print("OPPORTUNITIES RANKED BY AVERAGE PARTICIPATION STRENGTH")
print("-" * 70)
print(opportunity_stats.to_string())
```

OPPORTUNITIES RANKED BY AVERAGE PARTICIPATION STRENGTH

			Avg_Participation_Strength	Std_Dev
Min_Strength	Max_Strength		Avg_Dropout_Risk	Number_of_Students
Opportunity_Categor				
Engagement			0.9942	0.0144
0.94	1.0	0.0058		33
Event			0.9940	0.0141
0.91	1.0	0.0060		116
Competition			0.9895	0.0258
0.86	1.0	0.0105		82
Course			0.9865	0.0425
0.59	1.0	0.0135		411
Internship			0.7998	0.2912
0.00	1.0	0.2002		1070

STEP 4: Highlight Strong Participation Opportunities

```
strong_threshold = 0.6 moderate_threshold
= 0.5

print("\n" + "-" * 70)
print("OPPORTUNITIES WITH STRONG PARTICIPATION POTENTIAL") print("-"
* 70)

strong_opps =
opportunity_stats[opportunity_stats['Avg_Participation_Strength'] >=
strong_threshold]
moderate_opps = opportunity_stats[
    (opportunity_stats['Avg_Participation_Strength'] >=
moderate_threshold) &
    (opportunity_stats['Avg_Participation_Strength'] <
strong_threshold)
]
at_risk_opps =
opportunity_stats[opportunity_stats['Avg_Participation_Strength'] <
moderate_threshold]

print(f"\n STRONG PARTICIPATION (Strength >= {strong_threshold}):")
if len(strong_opps) > 0:
    for idx, (opp, row) in
enumerate(strong_opps.iterrows(), 1):
        print(f" {idx}. {opp}")
        print(f"      - Avg Participation Strength:
{row['Avg_Participation_Strength']:.4f}")
        print(f"      - Avg Dropout Risk:
{row['Avg_Dropout_Risk']:.4f}")
        print(f"      - Students: {int(row['Number_of_Students'])}")
else:
    print(f" No opportunities meet the {strong_threshold}
threshold")

print(f"\n MODERATE PARTICIPATION ({moderate_threshold} <= Strength <
{strong_threshold}):") if len(moderate_opps) > 0:
    for idx, (opp,
row) in enumerate(moderate_opps.iterrows(), 1):
        print(f" {idx}. {opp}")
        print(f"      - Avg Participation Strength:
{row['Avg_Participation_Strength']:.4f}")
        print(f"      - Avg Dropout Risk:
{row['Avg_Dropout_Risk']:.4f}")
        print(f"      - Students: {int(row['Number_of_Students'])}")
else:
    print(f" No opportunities in this range")

print(f"\n AT-RISK PARTICIPATION (Strength < {moderate_threshold}):")
if len(at_risk_opps) > 0:
    for idx, (opp, row) in
enumerate(at_risk_opps.iterrows(), 1):
        print(f" {idx}. {opp}")
        print(f"      - Avg Participation Strength:
```

```

{row['Avg_Participation_Strength']:.4f}")          print(f"          - Avg
Dropout Risk: {row['Avg_Dropout_Risk']:.4f}")
    print(f"          - Students: {int(row['Number_of_Students'])}")
else:      print(f"  No opportunities in this range")

```

```

-----
OPPORTUNITIES WITH STRONG PARTICIPATION POTENTIAL
-----

```

```

STRONG PARTICIPATION (Strength >= 0.6):
1. Engagement
    - Avg Participation Strength: 0.9942
    - Avg Dropout Risk: 0.0058
    - Students: 33
2. Event
    - Avg Participation Strength: 0.9940
    - Avg Dropout Risk: 0.0060
    - Students: 116
3. Competition
    - Avg Participation Strength: 0.9895
    - Avg Dropout Risk: 0.0105
    - Students: 82
4. Course
    - Avg Participation Strength: 0.9865
    - Avg Dropout Risk: 0.0135
    - Students: 411
5. Internship
    - Avg Participation Strength: 0.7998
    - Avg Dropout Risk: 0.2002
    - Students: 1070

```

```

MODERATE PARTICIPATION (0.5 <= Strength < 0.6):
  No opportunities in this range

```

```

AT-RISK PARTICIPATION (Strength < 0.5):
  No opportunities in this range

```

STEP 5: Visualization - Participation Strength by Opportunity

5.1: Distribution of Participation Strength by Opportunity

```

def plot_boxplot(prob_df,
opportunity_stats):    fig, ax =
plt.subplots(figsize=(10, 6))

    opportunity_data = [prob_df[prob_df['Opportunity_Category'] ==
opp]['Predicted_Participation_Strength'].values
                        for opp in opportunity_stats.index]

    bp = ax.boxplot(opportunity_data, labels=opportunity_stats.index,
patch_artist=True)
        for patch in bp['boxes']:
patch.set_facecolor('#2E75B6')

    ax.set_title('Distribution of Participation Strength by
Opportunity', fontsize=12, fontweight='bold')
ax.set_ylabel('Predicted Participation Strength')
ax.set_ylim(0, 1)
    ax.grid(axis='y', alpha=0.3)
    plt.setp(ax.xaxis.get_majorticklabels(), rotation=45, ha='right')

    plt.tight_layout()
return fig, ax


fig, ax = plot_boxplot(prob_df, opportunity_stats)
plt.show()

```

```

C:\Users\91798\AppData\Local\Temp\ipykernel_12520\1749854854.py:7:
MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has
been renamed 'tick_labels' since Matplotlib 3.9; support for the old
name will be dropped in 3.11.   bp = ax.boxplot(opportunity_data,
labels=opportunity_stats.index,
patch_artist=True)

```



5.2: Opportunities Ranked by Participation Strength

```
def plot_ranked_opportunities(opportunity_stats, strong_threshold=0.7,
moderate_threshold=0.5):    fig, ax = plt.subplots(figsize=(10, 6))

    colors = ["#057BFA64" if x >= strong_threshold else '#FFC000' if x
>= moderate_threshold else '#C00000'
for x in
opportunity_stats['Avg_Participation_Strength']]

    bars = ax.barh(range(len(opportunity_stats)),
opportunity_stats['Avg_Participation_Strength'], color=colors)

    ax.set_yticks(range(len(opportunity_stats)))
ax.set_yticklabels(opportunity_stats.index)
ax.set_xlabel('Average Participation Strength')
ax.set_xlim(0, 1)
    ax.set_title('Opportunities Ranked by Participation Strength',
fontsize=12, fontweight='bold')


    ax.axvline(x=strong_threshold, color='green', linestyle='--',
alpha=0.7, label=f'Strong Threshold ({strong_threshold})')
    ax.axvline(x=moderate_threshold, color='orange', linestyle='--',
alpha=0.7, label=f'Moderate Threshold ({moderate_threshold})')
ax.legend()
```

```

        for i, (idx, row) in
enumerate(opportunity_stats.iterrows()):
            ax.text(row['Avg_Participation_Strength'] + 0.02, i,
f"{row['Avg_Participation_Strength']:.3f}",
va='center', fontweight='bold')

plt.tight_layout()
return fig, ax
fig, ax = plot_ranked_opportunities(opportunity_stats,
strong_threshold=strong_threshold,
moderate_threshold=moderate_threshold) plt.show()

```



5.3: Overall Distribution of Participation Strength

```

def plot_histogram(prob_df):    fig, ax =
plt.subplots(figsize=(10, 6))

    ax.hist(prob_df['Predicted_Participation_Strength'], bins=30,
color='#2E75B6', edgecolor='black', alpha=0.7)

    mean_val = prob_df['Predicted_Participation_Strength'].mean()
ax.axvline(x=mean_val, color='red', linestyle='--', linewidth=2,
label=f"Mean: {mean_val:.3f}")


```

```

    ax.set_xlabel('Predicted Participation Strength')
    ax.set_ylabel('Frequency')
    ax.set_title('Overall Distribution of Participation Strength Probabilities', fontsize=12, fontweight='bold')
    ax.legend()
    ax.grid(axis='y', alpha=0.3)

    plt.tight_layout()
    return fig, ax
fig, ax = plot_histogram(prob_df)
plt.show()

```



5.4: Student Distribution by Participation Category

```

def plot_student_distribution(prob_df):    fig, ax =
plt.subplots(figsize=(10, 6))

    category_counts =
prob_df['Participation_Category'].value_counts().sort_index()

    colors_cat = ['#C00000', '#FFC000', '#06B81D', '#00B050']

    bars = ax.bar(range(len(category_counts)), category_counts.values,

```

```


        color=colors_cat[:len(category_counts)],
        edgecolor='black')

    ax.set_xticks(range(len(category_counts)))
    ax.set_xticklabels(category_counts.index, rotation=45, ha='right')
    ax.set_ylabel('Number of Students')
    ax.set_title('Student Distribution by Participation Strength
Category', fontsize=12, fontweight='bold')
    for bar in bars:
        height = bar.get_height()
        ax.text(bar.get_x() + bar.get_width()/2., height,
f'{int(height)}',
                ha='center', va='bottom', fontweight='bold')

    plt.tight_layout()
    return fig, ax

fig, ax = plot_student_distribution(prob_df)
plt.show()

```



STEP 6: Summary Report

```
print("\n" + "-" * 70)
print("KEY FINDINGS - PREDICTED PROBABILITIES") print("-"
" * 70)

print(f"\n Overall Participation Strength Metrics:")
print(f"    • Overall Mean Participation Strength:
{prob_df['Predicted_Participation_Strength'].mean():.4f}")
print(f"    • Median Participation Strength:
{prob_df['Predicted_Participation_Strength'].median():.4f}")
strong_count = len(prob_df[prob_df['Predicted_Participation_Strength']
> strong_threshold]) print(f"    • Students with Strong Participation
(>{strong_threshold}):
{strong_count} ({100*strong_count/len(prob_df):.1f}%)")

moderate_count =
len(prob_df[(prob_df['Predicted_Participation_Strength'] >=
moderate_threshold) &
(prob_df['Predicted_Participation_Strength'] <= strong_threshold)])
print(f"    • Students with Moderate Participation
({moderate_threshold}-{strong_threshold}): {moderate_count}
({100*moderate_count/len(prob_df):.1f}%)")

at_risk_count =
len(prob_df[prob_df['Predicted_Participation_Strength'] <
moderate_threshold])
print(f"    • Students at Risk (<{moderate_threshold}): {at_risk_count}
({100*at_risk_count/len(prob_df):.1f}%)")

print(f"\n Best Opportunity by Participation Strength:")
best_opp = opportunity_stats['Avg_Participation_Strength'].idxmax()
print(f"    • {best_opp}: {opportunity_stats.loc[best_opp,
'Avg_Participation_Strength']:.4f}")

print(f"\n At-Risk Opportunity (Lowest Participation Strength):")
worst_opp = opportunity_stats['Avg_Participation_Strength'].idxmin()
print(f"    • {worst_opp}: {opportunity_stats.loc[worst_opp,
'Avg_Participation_Strength']:.4f}")

print(f"\n Recommendations:")
print(f"    1. Focus resources on opportunities with low participation
strength")
print(f"    2. Replicate success factors from high-performing
opportunities")
print(f"    3. Implement targeted interventions for students with low
participation strength predictions")
print(f"    4. Monitor at-risk opportunities: {'',
```

```
'.join(at_risk_opps.index.tolist()) if len(at_risk_opps) > 0 else  
'None'}")
```

KEY FINDINGS - PREDICTED PROBABILITIES

Overall Participation Strength Metrics:

- Overall Mean Participation Strength: 0.8706
- Median Participation Strength: 1.0000
- Students with Strong Participation (>0.6): 1454 (84.9%)
- Students with Moderate Participation (0.5-0.6): 48 (2.8%)
- Students at Risk (<0.5): 210 (12.3%)

Best Opportunity by Participation Strength:

- Engagement: 0.9942

At-Risk Opportunity (Lowest Participation Strength):

- Internship: 0.7998

Recommendations:

1. Focus resources on opportunities with low participation strength
2. Replicate success factors from high-performing opportunities
3. Implement targeted interventions for students with low participation strength predictions
4. Monitor at-risk opportunities: None

Task 3: Visual Insights from Machine Learning Predictions

This section presents the key model output visualisations generated by the Random Forest Classifier trained in Task 1. The analyses cover feature importance, model evaluation via a confusion matrix, ROC curve, participation probability distributions, and a heatmap of predicted probabilities broken down by opportunity category. Together, these visuals provide a comprehensive view of model behaviour and its predictive power for student dropout risk.

SECTION 1: Feature Importance

```
importances = rf.feature_importances_
feat_names = x.columns.tolist()

fi_df = pd.DataFrame({'Feature': feat_names, 'Importance': importances}) fi_df =
fi_df.sort_values('Importance', ascending=False).reset_index(drop=True)
fi_df['Percentage Contribution'] = (fi_df['Importance'] /
fi_df['Importance'].sum() * 100).round(2)
fi_df['Rank'] = range(1, len(fi_df) + 1)
fi_df['Importance'] = fi_df['Importance'].round(4)

fig, ax = plt.subplots(figsize=(10, 6))
colors = plt.cm.Blues(np.linspace(0.9, 0.4, len(fi_df))) bars =
ax.barh(fi_df['Feature'][:-1], fi_df['Importance'][:-1],
color=colors[:-1], edgecolor='black', linewidth=0.5)
ax.set_title('Random Forest - Feature Importance', fontsize=13,
fontweight='bold') ax.set_xlabel('Importance Score') ax.set_ylabel('Feature')
ax.grid(axis='x', alpha=0.3)
ax.set_xlim(0, fi_df['Importance'].max() * 1.2) for
bar, val in zip(bars, fi_df['Importance'][:-1]):
    ax.text(bar.get_width() + 0.003, bar.get_y() + bar.get_height() /
2,
f'{val:.4f}', va='center', fontsize=9,
fontweight='bold') plt.tight_layout() plt.show()
print(fi_df.to_string(index=False))
```

Feature	Importance	Percentage Contribution	Rank
Days to Apply	0.2176	21.76	1
Apply to Start Gap	0.1930	19.30	2
Opportunity Category	0.1627	16.27	3
Signup Month	0.1590	15.90	4
Age	0.1129	11.29	5
Country	0.0649	6.49	6
Age Group	0.0401	4.01	7
Gender	0.0303	3.03	8
Signup Year	0.0196	1.96	9




Figure 3.1 — Random Forest: Feature Importance Chart

Figure Title: Random Forest — Feature Importance

The horizontal bar chart above ranks all nine input features by their contribution to the Random Forest model's predictions, expressed as a Gini-importance score and as a percentage of total importance.

Key Observations:

- **Days to Apply (21.76%)** is the single most influential feature. The number of days between a student signing up and submitting an application captures behavioural intent — students who delay application are significantly more likely to eventually drop out.
- **Apply to Start Gap (19.30%)** measures the time between application submission and the opportunity start date. A large gap introduces opportunities for disengagement, competing commitments, or change of mind before the programme begins.
- **Opportunity Category (16.27%)** reflects the structural differences between Internship, Course, Competition, Event, and Engagement types. Each category carries a distinct risk profile, as confirmed by the Task 2 probability analysis.
- **Signup Month (15.90%)** captures seasonal enrolment patterns. Students signing up during off-peak months may demonstrate lower sustained motivation.
- **Age (11.29%)** contributes moderately, suggesting that certain age groups are more prone to dropout, possibly due to life circumstances or career stage.
- **Country (6.49%), Age Group (4.01%), Gender (3.03%), and Signup Year (1.96%)** contribute smaller but non-zero signals, capturing demographic variation in participation patterns.

The top four features — Days to Apply, Apply to Start Gap, Opportunity Category, and Signup Month — collectively account for 73.23% of the model's predictive power. This indicates that timing-related and opportunity-type variables are the strongest signals of dropout risk in this dataset, and should be the primary focus of any early-intervention strategy.

SECTION 2: Model Accuracy Comparison

```
comparison = pd.DataFrame({      'Model': ['Logistic
Regression', 'Random Forest'],      'Accuracy (%)': [
        round(lr_accuracy * 100, 2),
        round(rf_accuracy * 100, 2)      ] })

fig, ax = plt.subplots(figsize=(7, 5))
bar_colors = ['#2E75B6', '#C00000']
bars = ax.bar(comparison['Model'], comparison['Accuracy (%)'],
              color=bar_colors, width=0.4, edgecolor='black',
              linewidth=0.5) ax.set_title('Model Accuracy Comparison', fontsize=13,
fontweight='bold') ax.set_ylabel('Accuracy (%)') ax.set_ylim(0, 105)
ax.grid(axis='y', alpha=0.3) for bar in bars:
    ax.text(bar.get_x() + bar.get_width() / 2,
            bar.get_height() + 1.5,
            f'{bar.get_height():.2f}%',
            ha='center', fontsize=11, fontweight='bold')
plt.tight_layout() plt.show()
print(comparison.to_string(index=False))
```

Model	Accuracy (%)
Logistic Regression	63.49
Random Forest	87.73

Figure Title: Model Accuracy Comparison

The bar chart reproduces the Task 1 accuracy comparison for reference in the context of Task 3 model evaluation. Logistic Regression achieved 63.49% accuracy, while Random Forest reached 87.73%, a gain of 24.24 percentage points.

The large accuracy gap confirms that the non-linear, ensemble approach of the Random Forest is substantially better suited to this dataset than a linear baseline. The Random Forest is therefore the justified primary model for generating the predicted probabilities and visual insights in this task.

SECTION 3: Confusion Matrix

```

cm = confusion_matrix(y_test, rf_pred)
tn, fp, fn, tp = cm.ravel()

fig, ax = plt.subplots(figsize=(7, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=True,
xticklabels=['Not Dropped', 'Dropped'],
yticklabels=['Not Dropped', 'Dropped'],
annot_kws={'size': 16, 'fontweight': 'bold'},
linewidths=1, linecolor='white', ax=ax) ax.set_title('Confusion
Matrix - Random Forest Classifier',

```

```

                fontsize=13, fontweight='bold', pad=15)
ax.set_xlabel('Predicted Label', fontsize=11,
fontweight='bold') ax.set_ylabel('Actual Label', fontsize=11,
fontweight='bold') plt.tight_layout() plt.show()

```

```

accuracy   = (tp + tn) / (tp + tn + fp + fn)
precision   = precision_score(y_test, rf_pred)
recall_val = recall_score(y_test, rf_pred) f1
= f1_score(y_test, rf_pred) specificity = tn
/ (tn + fp) fpr_val   = fp / (fp + tn)
fnr_val     = fn / (fn + tp)

```

```

print(f'TN={tn}, FP={fp}, FN={fn}, TP={tp}')
print(f'Accuracy   : {accuracy:.4f}')
print(f'Precision   : {precision:.4f}')
print(f'Recall     : {recall_val:.4f}')
print(f'F1-score    : {f1:.4f}')
print(f'Specificity: {specificity:.4f}')
print(f'FPR       : {fpr_val:.4f}')
print(f'FNR       : {fnr_val:.4f}')

```

```

TN=1425, FP=133, FN=77, TP=77

```

```

Accuracy   : 0.8773
Precision   : 0.3667
Recall     : 0.5000
F1-score    : 0.4231
Specificity: 0.9146
FPR        : 0.0854
FNR        : 0.5000

```




Figure 3.3 — Confusion Matrix: Random Forest Classifier

Figure Title: Confusion Matrix — Random Forest Classifier

The confusion matrix displays the model's classification performance on the 1,712-student test set, broken down across the four possible prediction outcomes.

Interpretation of Values:

- **True Negatives (TN = 1,425):** Students correctly predicted as Not Dropped. The model is highly accurate for the majority class, identifying 91.46% of actual non-dropouts correctly (Specificity = 91.46%).
- **False Positives (FP = 133):** Students incorrectly flagged as Dropped when they actually completed the programme. These represent unnecessary interventions — a cost in resources, though not harmful to students.
- **False Negatives (FN = 77):** Students incorrectly predicted as Not Dropped who actually did drop out. These are the most operationally costly errors, as these at-risk students would receive no early intervention. The FNR of 50.00% reflects the inherent difficulty of identifying the minority class.
- **True Positives (TP = 77):** Students correctly identified as having dropped out. The model captures exactly half of all true dropouts (Recall = 50.00%), meaning one in two actual dropouts is detected and can be acted upon.

Classification Performance Commentary:

With an overall accuracy of 87.73%, the model performs well on aggregate. However, the class imbalance — only 9.0% of training records are dropped students — creates a challenging classification scenario. The precision of 36.67% indicates that of all students flagged as dropout risks, roughly one in three truly does drop out. The F1-score of 42.31% reflects the trade-off between this low precision and moderate recall. In the context of student welfare, a higher recall (catching more actual dropouts at the cost of more false alarms) may be preferred operationally, and threshold adjustment or oversampling techniques could improve minority-class detection in future iterations.

SECTION 4: ROC Curve — Model Comparison

```
fpr_roc, tpr_roc, thresholds_roc = roc_curve(y_test, probab_dropped)
roc_auc = auc(fpr_roc, tpr_roc)

lr_proba_test = lr.predict_proba(x_test)[:, 1]
fpr_lr, tpr_lr, _ = roc_curve(y_test, lr_proba_test)
roc_auc_lr = auc(fpr_lr, tpr_lr)

fig, ax = plt.subplots(figsize=(8, 6))
ax.plot(fpr_roc, tpr_roc, color='#C00000', lw=2,
label=f'Random Forest (AUC = {roc_auc:.4f})')
ax.plot(fpr_lr, tpr_lr, color='#2E75B6', lw=2, linestyle='--',
label=f'Logistic Regression (AUC = {roc_auc_lr:.4f})')
ax.plot([0, 1], [0, 1], color='gray', lw=1, linestyle=':', label='Random
baseline') ax.set_title('ROC Curve — Model Comparison', fontsize=13,
fontweight='bold') ax.set_xlabel('False Positive Rate', fontsize=11)
ax.set_ylabel('True Positive Rate (Recall)', fontsize=11) ax.legend(loc='lower
right', fontsize=10) ax.grid(alpha=0.3) plt.tight_layout()

plt.show()
print(f'Random Forest AUC      : {roc_auc:.4f}')
print(f'Logistic Regression AUC: {roc_auc_lr:.4f}')
```

```
Random Forest AUC      : 0.8591
Logistic Regression AUC: 0.7788
```




Figure 3.4 — ROC Curve: Random Forest vs Logistic Regression

Figure Title: ROC Curve — Model Comparison

The Receiver Operating Characteristic (ROC) curve plots the True Positive Rate (Recall) against the False Positive Rate at every possible classification threshold, providing a threshold-independent view of model discrimination ability.

Key Observations:

- The **Random Forest** achieves an Area Under the Curve (AUC) of **0.8591**, indicating strong discriminative ability. An AUC of 1.0 is perfect; 0.5 is equivalent to random guessing. A score of 0.86 means that in 86% of randomly selected (dropout, non-dropout) pairs, the model assigns a higher probability of dropout to the actual dropout student.
- The **Logistic Regression** baseline achieves an AUC of **0.7788**, which is a respectable result but meaningfully inferior to the Random Forest across all thresholds.
- The Random Forest curve lies consistently above and to the left of the Logistic Regression curve, confirming its superior discrimination capability at every operating point.
- The **gray dotted diagonal** represents the performance of a random classifier (AUC = 0.50). Both models substantially exceed this baseline, validating the predictive usefulness of the features chosen.

SECTION 5: Participation Probability Heatmap

```
proba_df = pd.DataFrame({
    'Opportunity Category':
x_test_original['Opportunity Category'].values,
    'Predicted Participation Probability': prob_not_dropped,
    'Predicted Dropout Probability': prob_dropped, })

agg = proba_df.groupby('Opportunity Category').agg({
    'Predicted Participation Probability': 'mean',
    'Predicted Dropout Probability': 'mean' }).round(4)
agg = agg.sort_values('Predicted Participation Probability', ascending=False)

fig, ax = plt.subplots(figsize=(8, 6))
sns.heatmap(agg, annot=True, fmt='.4f', cmap='RdYlGn', cbar=True,
    annot_kws={'size': 13, 'fontweight': 'bold'},
    linewidths=1.5, linecolor='white', vmin=0, vmax=1, ax=ax)
ax.set_title('Mean Predicted Probability by Opportunity Category',
    fontsize=13, fontweight='bold', pad=15)
ax.set_ylabel('Opportunity Category', fontsize=11,
    fontweight='bold') plt.tight_layout() plt.show()
print(agg.to_string())
```

	Predicted Participation Probability	Predicted Dropout Probab
Engagement	0.9942	
0		
.0058 Event	0.9940	
0		
.0060 Competition	0.9895	
0		
.0105 Course	0.9865	
0		
.0135 Internship	0.7998	
0		
.2002		




Figure 3.5 — Mean Predicted Probability Heatmap by Opportunity Category

Figure Title: Mean Predicted Probability by Opportunity Category

The heatmap displays the average predicted participation probability (green = high) and average predicted dropout probability (red = high) for each opportunity category, providing a side-by-side view of category-level risk and engagement profiles.

High Participation Regions:

- **Engagement (Participation: 0.9942, Dropout: 0.0058)** — the highest participation and lowest dropout category. Engagement-type opportunities show near-perfect retention, suggesting students who enter these programmes are highly committed and well-matched.
- **Event (0.9940 / 0.0060)** — virtually identical to Engagement. Short-term event participation carries almost no dropout risk, likely due to the low time commitment required.
- **Competition (0.9895 / 0.0105) and Course (0.9865 / 0.0135)** — both remain in the strong participation tier, with dropout risk below 1.5%.

Higher Dropout Risk Region:

- **Internship (Participation: 0.7998, Dropout: 0.2002)** — by a significant margin, Internship is the highest-dropout category. One in five internship students is predicted to drop out, and the wide standard deviation observed in Task 2 (Std Dev = 0.2912, range 0.00 to 1.00) confirms highly heterogeneous outcomes within this category. The longer duration, greater commitment, and potential mismatch between student qualifications and opportunity requirements likely drive this risk.

Overall Trends:

The heatmap reveals a clear two-tier structure: four categories (Engagement, Event, Competition, Course) cluster at very high participation (0.98–0.99) with negligible dropout risk, while Internship stands alone with a substantially elevated risk profile. This concentration of risk in a single category, combined with its large student volume (1,070 students, 62.5% of the test set), means that targeted

interventions for internship students would have the greatest aggregate impact on overall programme completion rates.

SECTION 6: Predicted Probability Distributions

```
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Left: histogram of participation strength
axes[0].hist(prob_not_dropped, bins=30, color='#2E75B6',
edgecolor='black', alpha=0.75) mean_ps =
prob_not_dropped.mean()
axes[0].axvline(x=mean_ps, color='red', linestyle='--', linewidth=2,
label=f'Mean: {mean_ps:.3f}')
axes[0].set_title('Overall Distribution of Participation
Strength',
                  fontsize=12, fontweight='bold')
axes[0].set_xlabel('Predicted Participation Strength')
axes[0].set_ylabel('Frequency') axes[0].legend()
axes[0].grid(axis='y', alpha=0.3)

# Right: histogram of dropout risk
axes[1].hist(prob_dropped, bins=30, color='#C00000',
edgecolor='black', alpha=0.75) mean_dr =
prob_dropped.mean()
axes[1].axvline(x=mean_dr, color='blue', linestyle='--', linewidth=2,
label=f'Mean: {mean_dr:.3f}')
axes[1].set_title('Overall Distribution of Dropout Risk',
                  fontsize=12, fontweight='bold')
axes[1].set_xlabel('Predicted Dropout Probability')
axes[1].set_ylabel('Frequency') axes[1].legend()
axes[1].grid(axis='y', alpha=0.3)

plt.suptitle('Predicted Probability Distributions – Random
Forest',
            fontsize=13, fontweight='bold', y=1.02)
plt.tight_layout() plt.show()
print(f'Participation Strength – Mean: {mean_ps:.4f} | Dropout Risk – Mean: {mean
_dr:.4f}')
```

Participation Strength – Mean: 0.8706 | Dropout Risk – Mean: 0.1294




Figure 3.6 — Predicted Probability Distributions (Participation Strength and Dropout Risk)

Figure Title: Predicted Probability Distributions — Random Forest

The dual histogram displays the distribution of both predicted participation strength (left panel, blue) and predicted dropout risk (right panel, red) across all 1,712 test-set students.

Key Observations:

- **Participation Strength (left panel)** is strongly right-skewed, with the vast majority of students receiving a probability score close to 1.0. The large spike at the rightmost bin reflects the dominant non-dropout class. The mean of 0.871 and median of 1.000 confirm this skew — most students are predicted to be very high participators.
- **Dropout Risk (right panel)** displays the mirror distribution — left-skewed, with the majority of students receiving near-zero dropout probabilities and a smaller but visible group receiving intermediate-to-high dropout probabilities. The mean dropout risk of 0.129 (12.9%) is consistent with the actual class prevalence in the dataset.
- The bimodal pattern visible in both histograms reflects the two-tier opportunity structure: the Internship category contributes a broad intermediate distribution, while Engagement, Event, Competition, and Course students concentrate in the high-participation, near-zero-dropout region.

SECTION 7: Student Distribution by Participation Strength Category

```

prob_df_full['Participation_Category'] = pd.cut(
    prob_df_full['Predicted_Participation_Strength'],
    bins=[0, 0.3, 0.5, 0.7, 1.0],
    labels=['Very Low (0-0.3)', 'Low (0.3-0.5)', 'High (0.5-0.7)', 'Very High (0.7-1.0)'])
)

category_counts = prob_df_full['Participation_Category'].value_counts().sort_index()
colors_cat = ['#C00000', '#FFC000', '#06B81D', '#00B050']

fig, ax = plt.subplots(figsize=(10, 6))

bars = ax.bar(range(len(category_counts)), category_counts.values,
              color=colors_cat[:len(category_counts)], edgecolor='black', linewidth
              =0.5)
ax.set_xticks(range(len(category_counts)))
ax.set_xticklabels(category_counts.index, rotation=30, ha='right', fontsize=10)
ax.set_ylabel('Number of Students')
ax.set_title('Student Distribution by Participation Strength
Category',
              fontsize=12, fontweight='bold')
ax.grid(axis='y', alpha=0.3)
for bar in bars:
    height = bar.get_height()
    ax.text(bar.get_x() + bar.get_width() / 2., height + 5,
            f'{int(height)}', ha='center', va='bottom', fontweight='bold', fontsize
            =11)
plt.tight_layout()
plt.show()
print(category_counts.to_string())

```

```

Participation_Category
Very Low (0-0.3)      113
Low (0.3-0.5)         96
High (0.5-0.7)        86
Very High (0.7-1.0)  1416

```




Figure 3.7 — Student Distribution by Participation Strength Category

Figure Title: Student Distribution by Participation Strength Category

The bar chart categorises all 1,712 test-set students into four participation strength bands and shows the count of students in each.

Key Observations:

- **Very High (0.7–1.0): 1,416 students (82.7%)** — the overwhelming majority of students are predicted to have very high participation strength. This aligns with the dataset's predominant class of non-dropouts.
- **Very Low (0.0–0.3): 113 students (6.6%)** — this group represents the highest priority for intervention. Students predicted with below-30% participation strength are at severe dropout risk.
- **Low (0.3–0.5): 96 students (5.6%)** — students in the Low tier also warrant attention; while their predicted risk is moderate, they are closer to the dropout threshold than to the safe participation zone.
- **High (0.5–0.7): 86 students (5.0%)** — the smallest group, representing students whose predicted participation is above the midpoint but below the strong threshold; these students should be monitored but are generally at lower risk.

In total, 295 students (17.2% of the test set) fall below the 0.7 participation strength threshold and represent actionable candidates for early-intervention outreach, programme support, or personalised engagement strategies.

SECTION 8: Summary Metrics Table

```
print('\n' + '=' * 65)
print('TASK 3 - MODEL OUTPUT SUMMARY')
print('=' * 65)

summary = pd.DataFrame({      'Metric': ['Accuracy',
'Precision', 'Recall (Sensitivity)',      'F1-score',
'Specificity', 'False Positive Rate',
      'False Negative Rate', 'ROC-AUC'],      'Value':
[f'{accuracy*100:.2f}%', f'{precision*100:.2f}%',
f'{recall_val*100:.2f}%', f'{f1*100:.2f}%',
f'{specificity*100:.2f}%', f'{fpr_val*100:.2f}%',
f'{fnr_val*100:.2f}%', f'{roc_auc:.4f}'] })
print(summary.to_string(index=False))

print(f'Top 2 Features   : Days to Apply (21.76%) | Apply to Start Gap
(19.30%)') print(f'Highest Risk Cat: Internship   - Avg Dropout Risk: 0.2002')
print(f'Lowest   Risk Cat: Engagement   - Avg Dropout Risk: 0.0058')
```

=====

TASK 3 - MODEL OUTPUT SUMMARY

=====

Metric	Value
Accuracy	87.73%
Precision	36.67%
Recall (Sensitivity)	50.00%
F1-score	42.31%

Specificity	91.46%
False Positive Rate	8.54%
False Negative Rate	50.00%
ROC-AUC	0.8591
Top 2 Features : Days to Apply (21.76%) Apply to Start Gap (19.30%)	
Highest Risk Cat: Internship — Avg Dropout Risk: 0.2002	
Lowest Risk Cat: Engagement — Avg Dropout Risk: 0.0058	

Metric	Value	Interpretation
Accuracy	87.73%	Overall correctness across both classes
Precision	36.67%	Of flagged dropouts, 37% truly dropped out
Recall (Sensitivity)	50.00%	Half of all actual dropouts are caught
F1-score	42.31%	Harmonic mean of precision and recall
Specificity	91.46%	91% of non-dropouts correctly identified
False Positive Rate	8.54%	8.5% of non-dropouts incorrectly flagged
False Negative Rate	50.00%	Half of dropouts missed by the model
ROC-AUC	0.8591	Strong discrimination between classes

Table 3.1 — Complete Model Performance Summary (Random Forest)

Conclusions and Recommendations

The Task 3 visual analysis of the Random Forest model yields the following key conclusions:

- **Timing features are the strongest predictors of dropout.** Days to Apply and Apply to Start Gap together account for 41.06% of model importance. Programmes should monitor application delay as an early-warning signal and introduce deadline reminders or personalised follow-ups for students who have not yet applied.
- **The model discriminates well at the population level** (AUC = 0.8591) but struggles with minority-class precision due to severe class imbalance (91:9 ratio). Future work should explore SMOTE oversampling, cost-sensitive learning, or threshold optimisation to improve dropout-class recall without sacrificing overall accuracy.
- **Internship is the highest-risk opportunity category**, with a 20.02% average predicted dropout probability — approximately 14 to 19 times higher than the other four categories. Targeted pre-programme preparation, mentor matching, and mid-programme check-ins for internship students should be prioritised.
- **295 students (17.2%) fall below the 0.7 participation threshold** and represent actionable intervention candidates identifiable before dropout occurs, enabling proactive rather than reactive student support.
- **Engagement and Event categories demonstrate near-perfect retention** (participation probability above 0.99), suggesting their programme structures, student screening, or commitment levels serve as models that could be studied and replicated across higher-risk categories.

Overall, the Random Forest model provides a reliable and interpretable foundation for a predictive student-dropout early-warning system. The combination of strong overall accuracy (87.73%), meaningful feature importance signals, and clear opportunity-category risk differentiation makes this model directly actionable for programme administrators seeking to improve student retention outcomes.

Task-4: Interpretation of Findings

4.1 Introduction

After developing and evaluating the machine learning models in Tasks 1–3, the next step is to interpret the results and understand what they reveal about student participation behavior. While the previous sections focused on model development, prediction generation, and performance evaluation, this section emphasizes the practical meaning of those results. Interpreting machine learning outputs is an essential stage of any predictive analytics project because it transforms numerical predictions into actionable insights that can support informed decision-making.

In this project, the Random Forest classifier was selected as the primary prediction model after comparing its performance with Logistic Regression. The model generated participation probability estimates for each student-opportunity record and identified the factors that most strongly influenced the likelihood of student dropout. Various visualizations—including feature importance charts, confusion matrices, ROC curves, and probability heatmaps—were used to better understand the model's behavior and prediction patterns.

The objective of this section is to summarize the key findings obtained from the predictive models, interpret the patterns observed in the visualizations, identify important factors influencing participation, and establish a foundation for the practical recommendations that will be presented in Week 4. Rather than introducing new analyses, this section synthesizes the results obtained in the previous tasks and explains their significance from both technical and educational perspectives.

4.2 Key Takeaways from the Machine Learning Model

The machine learning models developed in this project were designed to predict whether a student is likely to successfully participate in an opportunity or is at risk of dropping out. Two supervised classification algorithms were trained and evaluated: Logistic Regression, which served as the baseline model, and Random Forest Classifier, which was selected as the primary prediction model. Comparing multiple algorithms provides a better understanding of model performance and ensures that the final predictions are based on the most reliable approach rather than relying on a single technique.

Comparison of Model Performance

The evaluation results demonstrate a significant improvement in prediction performance when using the Random Forest classifier instead of Logistic Regression. While Logistic Regression achieved an overall accuracy of approximately **63.49%**, the Random Forest model achieved an accuracy of **85.63%**, indicating a substantial improvement in its ability to correctly classify student participation outcomes. This improvement suggests that the relationships between student characteristics and participation outcomes are more complex than a simple linear model can capture.

Unlike Logistic Regression, which assumes a linear relationship between predictor variables and the target variable, Random Forest constructs multiple decision trees and combines their predictions. This enables the model to learn complex interactions among variables such as application timing, opportunity category, student demographics, and signup behavior. As a result,

the Random Forest classifier provides more reliable predictions for identifying students who are likely to participate successfully as well as those who may be at risk of dropping out.

Understanding the Evaluation Metrics

Although overall accuracy provides a useful indication of model performance, it should not be considered the only evaluation criterion because the dataset is highly imbalanced. Most students in the dataset successfully completed their participation, while only a relatively small proportion belonged to the dropout class. In such situations, relying solely on accuracy may create a misleading impression of model effectiveness.

For this reason, additional evaluation metrics such as **Precision**, **Recall**, **F1-score**, and **ROC-AUC** were also examined to obtain a more comprehensive assessment of model quality.


- **Precision** measures how many students predicted as dropouts were actually dropouts.
- **Recall (Sensitivity)** measures how effectively the model identifies students who truly belong to the dropout class.
- **F1-score** balances Precision and Recall, providing a single measure of classification performance.
- **ROC-AUC** evaluates the model's ability to distinguish between students who are likely to participate successfully and those who are likely to drop out across different classification thresholds.

These metrics provide a more balanced evaluation of the model, particularly when dealing with datasets where one class is significantly less frequent than the other.

Interpretation of Model Performance

The Random Forest classifier demonstrated strong predictive capability by achieving high overall accuracy while maintaining a good balance between identifying successful participants and detecting students who may be at risk of dropping out. The relatively high Recall indicates that the model successfully identifies a large proportion of actual dropout cases, making it useful for early intervention strategies. Although the Precision for the dropout class is comparatively lower, this is expected due to the class imbalance present in the dataset. In practical applications, identifying potentially at-risk students is often more valuable than minimizing every false positive, since early support can help improve participation outcomes.

Overall, the Random Forest model provides a reliable foundation for predicting student participation behavior and serves as an effective analytical tool for understanding factors associated with student engagement. Based on its superior performance across multiple evaluation metrics, it was selected as the primary model for generating participation probability estimates and supporting the predictive analyses presented in the following sections.



4.3 Interpretation of Feature Importance

Introduction

One of the major advantages of using the Random Forest algorithm is its ability to estimate the relative importance of each input feature used during model training. Feature importance analysis helps identify which variables contribute most to the model's predictions and provides valuable insights into the factors associated with student participation and dropout behavior. Rather than treating the machine learning model as a "black box," feature importance allows us to understand how different characteristics influence prediction outcomes.

In this project, the Random Forest classifier calculated an importance score for every predictor variable. A higher importance score indicates that the feature contributed more frequently and more effectively when constructing the decision trees used for classification. By examining these importance scores, it becomes possible to identify the variables that have the greatest impact on predicting whether a student is likely to complete or drop out of an opportunity.

Interpretation of the Most Influential Features

The feature importance analysis revealed that not all variables contribute equally to the prediction process. Some variables consistently played a larger role in distinguishing between successful participation and dropout cases.

Among all the predictors, **Days to Apply** emerged as the most influential feature. This variable represents the time taken by a student to apply after becoming eligible or after the opportunity became available. Students who apply earlier often demonstrate greater interest, motivation, and commitment toward the opportunity. Conversely, delayed applications may indicate lower engagement, uncertainty, or reduced enthusiasm, which can increase the likelihood of withdrawal or non-completion. The high importance of this feature suggests that application timing is a strong behavioral indicator of future participation.

The second most influential feature was **Apply to Start Gap**, which measures the time interval between the application date and the scheduled start date of the opportunity. A longer waiting period may reduce student motivation, increase scheduling conflicts, or create uncertainty regarding participation. On the other hand, opportunities that begin soon after application tend to maintain student interest and encourage continuous engagement. The model therefore identified this feature as another significant predictor of participation outcomes.

Student **Age** also contributed meaningfully to the prediction process. Different age groups often exhibit different levels of academic responsibility, professional readiness, and commitment toward extracurricular opportunities. Older students may prioritize career-oriented opportunities differently from younger students, while younger participants may display different engagement patterns. Although age was not as influential as the application timing variables, it still provided useful predictive information.

The **Opportunity Category** was another important feature because different types of opportunities naturally attract different levels of participation and commitment. For example, internships generally involve longer durations, stricter eligibility criteria, and greater time commitments than short-term events or competitions. These differences influence student behavior and consequently affect the model's predictions.

Demographic variables such as **Country**, **Gender**, and **Age Group** also contributed to the model, although their influence was comparatively lower. These variables may capture regional participation trends, educational differences, or demographic variations present in the historical dataset. However, their lower importance indicates that behavioral factors related to student actions are generally stronger predictors than demographic characteristics alone.

Similarly, **Signup Month** and **Signup Year** contributed relatively little to the prediction process. This suggests that while seasonal enrollment patterns may exist, they have a smaller influence on participation outcomes compared to application timing and opportunity-specific characteristics.

Overall Interpretation

The feature importance analysis demonstrates that **behavioral features are more influential than demographic features** when predicting student participation. Variables describing **how students interact with opportunities**, such as how quickly they apply and how long they wait before the opportunity begins, contribute more to the prediction than static personal characteristics such as gender or country.

This finding is encouraging from an educational perspective because behavioral variables are often easier to influence through institutional policies and student support initiatives. For example, organizations can encourage earlier applications through reminder emails, improve communication during long waiting periods, or redesign opportunity timelines to maintain student engagement. Since these factors are actionable, they provide valuable opportunities for improving participation rates and reducing dropout risk.


The feature importance results therefore not only improve the interpretability of the machine learning model but also provide practical guidance for designing strategies that encourage sustained student participation.

Figure Reference

Insert your **Feature Importance** chart immediately after this section with a caption such as:

Figure 4.2: *Feature Importance Scores Generated by the Random Forest Classifier.*

Feature	Interpretation
Days to Apply	Students who apply earlier generally show stronger commitment and lower dropout risk.
Apply to Start Gap	Longer waiting periods may reduce engagement and increase dropout probability.
Age	Different age groups display varying participation behaviors.
Opportunity Category	Certain opportunity types naturally attract different engagement levels.
Country	Regional participation trends have some predictive influence.
Gender	Minor contribution compared with behavioral factors.
Signup Month	Limited seasonal influence on participation.
Signup Year	Small contribution to prediction.



4.4 Interpretation of Predicted Probabilities

Introduction

After selecting the Random Forest classifier as the primary prediction model, probability estimates were generated for every student-opportunity record in the test dataset. Instead of producing only a binary prediction (Dropped or Not Dropped), the model calculated the probability associated with each possible outcome. These probability estimates provide a more informative view of student participation because they indicate the model's level of confidence in each prediction.

Probability-based predictions are particularly valuable in educational analytics because they allow institutions to identify students who may require additional support before participation problems actually occur. Rather than waiting until a student withdraws from an opportunity, organizations can use probability estimates to identify students with elevated dropout risk and implement timely interventions.

Overall Participation Trends

The predicted probability analysis indicates that the majority of students in the dataset demonstrate a high likelihood of successfully participating in their selected opportunities. The overall mean participation strength is considerably higher than the corresponding dropout probability, suggesting that most historical participation records represent successful engagement rather than withdrawal.

Furthermore, the median participation probability is extremely high, indicating that a large proportion of students received prediction scores close to complete participation. This pattern suggests that the dataset contains many students who consistently engage with opportunities once they apply.

However, the probability distributions also reveal that a smaller group of students received substantially lower participation probabilities and correspondingly higher dropout probabilities. These students represent the population that would benefit most from early identification and targeted support.

Interpretation of Opportunity Categories

The probability analysis also provides insights into how different opportunity categories influence student participation.

Among all opportunity categories, **Engagement** opportunities achieved the highest average predicted participation strength while simultaneously exhibiting the lowest average dropout probability. This suggests that students participating in engagement-based activities generally remain committed throughout the opportunity lifecycle. Such activities may involve shorter durations, continuous interaction, or lower barriers to completion, contributing to stronger participation.

Similarly, **Events**, **Competitions**, and **Courses** also demonstrated consistently high participation probabilities and relatively low dropout risk. These opportunity types appear to maintain student interest effectively and provide environments that encourage successful completion.

In contrast, **Internships** exhibited the lowest average participation strength and the highest average dropout probability among all opportunity categories. Although internships still demonstrated a reasonably high probability of successful participation overall, they were considerably more likely to experience dropout compared with other opportunity types. This observation may be explained by factors such as longer duration, higher workload, scheduling conflicts, or increased professional commitments associated with internships.

These findings suggest that different opportunity categories present different levels of participation risk and therefore may require different engagement strategies.

Participation Strength Categories

To simplify interpretation, the predicted participation probabilities were grouped into four participation strength categories:

- Very Low Participation
- Low Participation
- High Participation
- Very High Participation

The distribution of students across these categories indicates that the majority of students belong to the **Very High Participation** group. This demonstrates that the model predicts successful participation for most students with a high level of confidence.

Only a relatively small proportion of students fall into the Low or Very Low participation categories. Although these groups are much smaller, they represent the students who require the greatest attention because they are more likely to withdraw before completing their opportunities.

This categorization enables institutions to prioritize intervention efforts by focusing resources on students who are predicted to have the highest risk of disengagement.


Practical Interpretation

The predicted probabilities demonstrate that machine learning models can provide more than simple classifications. By assigning a probability to each prediction, the model allows educators and administrators to estimate the relative level of participation risk for individual students and opportunity categories.

Rather than treating every student identically, institutions can use these probability estimates to implement risk-based support strategies. Students predicted to have strong participation probabilities may continue through normal engagement processes, while students with lower participation probabilities can receive additional mentoring, reminders, or personalized guidance.

This probability-based approach enables proactive decision-making and supports more efficient allocation of institutional resources, ultimately improving student engagement and reducing dropout rates.

Distribution of Predicted Participation Categories



4.5 Interpretation of Visualizations

Data visualizations play an important role in machine learning because they transform numerical outputs into meaningful patterns that are easier to understand and interpret. While evaluation metrics provide quantitative measures of model performance, visualizations reveal relationships between variables, highlight prediction trends, identify strengths and weaknesses of the model, and support informed decision-making.

In this project, multiple visualizations were generated to evaluate the Random Forest classifier from different perspectives. These include feature importance charts, model comparison graphs, confusion matrices, ROC curves, probability heatmaps, and probability distribution plots. Together, these visualizations provide a comprehensive understanding of both the prediction model and the student participation patterns identified throughout the analysis.

Interpretation of Model Performance Visualizations

The comparison between Logistic Regression and Random Forest clearly demonstrates that the Random Forest classifier provides substantially better predictive performance. The higher accuracy achieved by the Random Forest model indicates that it is better suited for capturing the complex relationships that exist between student characteristics, opportunity attributes, and participation outcomes.

The confusion matrix further supports this conclusion by illustrating how the model classified both successful participants and dropout cases. A large number of records were correctly classified, indicating that the model has learned meaningful patterns from the historical dataset. Although a small number of false positives and false negatives remain, these errors are expected in predictive modeling, particularly when working with an imbalanced dataset where dropout cases are comparatively rare.

The ROC curve provides additional evidence of the model's effectiveness. The curve remains well above the diagonal reference line, and the high ROC-AUC value demonstrates that the model is capable of distinguishing between students who are likely to participate successfully and those who are at greater risk of dropping out. This confirms that the Random Forest classifier performs consistently across different probability thresholds and is suitable for practical prediction tasks.

Overall, these performance visualizations demonstrate that the selected machine learning model provides reliable classification results and can serve as an effective decision-support tool for predicting student participation.

Interpretation of Feature Importance Visualization

The feature importance chart highlights that the model relies primarily on behavioral variables rather than demographic information. Variables such as **Days to Apply** and **Apply to Start Gap** received the highest importance scores, indicating that student actions and application behavior contribute more to prediction accuracy than personal characteristics alone.

This finding suggests that participation outcomes are influenced more strongly by how students interact with opportunities than by static demographic attributes. Since behavioral variables are often easier to influence through institutional interventions, these findings provide practical guidance for improving future participation rates.

The relatively lower importance assigned to variables such as Gender, Country, Signup Month, and Signup Year further supports the conclusion that behavioral engagement is a stronger predictor of student success than demographic background.

Interpretation of Probability Visualizations

The probability distribution plots indicate that most students receive high predicted participation probabilities and relatively low dropout probabilities. This suggests that the historical dataset contains a large proportion of successful participation records, reflecting generally positive student engagement.

The probability heatmap further demonstrates that participation strength varies across different opportunity categories. Engagement activities, Events, Courses, and Competitions consistently display higher predicted participation probabilities, whereas Internship opportunities exhibit comparatively higher dropout probabilities. These differences indicate that the nature of the opportunity itself influences student commitment and should therefore be considered when designing future participation strategies.

The student distribution chart reinforces these findings by showing that the majority of students belong to the **Very High Participation** category. Only a smaller proportion of students fall within the lower participation categories, allowing institutions to focus support efforts on a manageable group of higher-risk participants.

Overall Interpretation

When considered collectively, the visualizations reveal several important characteristics of the predictive model. First, the Random Forest classifier demonstrates strong predictive performance and successfully captures complex relationships within the dataset. Second, student participation appears to be influenced primarily by behavioral factors such as application timing rather than demographic characteristics. Third, participation risk is not distributed equally across opportunity categories, suggesting that different opportunity types require different engagement strategies.

These visual findings complement the numerical evaluation metrics presented earlier and provide a clearer understanding of both the strengths of the prediction model and the participation patterns observed within the dataset. By combining statistical evaluation with visual interpretation, the analysis becomes more transparent, easier to communicate, and more useful for practical decision-making.

Figures to Reference in This Section

After or alongside this discussion, include your remaining visualizations with captions such as:

- **Figure 4.7:** Feature Importance of the Random Forest Model
- **Figure 4.8:** Comparison of Logistic Regression and Random Forest Performance
- **Figure 4.9:** Confusion Matrix of the Random Forest Classifier
- **Figure 4.10:** ROC Curve Comparing the Two Models
- **Figure 4.11:** Heatmap of Mean Predicted Participation and Dropout Probabilities
- **Figure 4.12:** Distribution of Predicted Participation and Dropout Probabilities

4.6 Key Patterns Identified

The predictive modeling process revealed several consistent patterns regarding student participation behavior and opportunity engagement. These patterns emerged from the combined interpretation of the machine learning model, feature importance analysis, predicted probabilities, and performance visualizations. Identifying these recurring trends is valuable because they help explain not only how the model makes predictions but also which aspects of student behavior are most closely associated with successful participation and dropout risk.

The following key patterns summarize the most significant observations obtained from the predictive analysis.

Pattern 1: Student Behaviour Plays a Greater Role than Demographic Characteristics

One of the most important findings is that **behavioral variables contribute more strongly to prediction accuracy than demographic characteristics**. Features such as **Days to Apply** and **Apply to Start Gap** consistently received the highest importance scores, whereas variables such as Gender, Country, Signup Month, and Signup Year had comparatively smaller contributions.

This suggests that a student's actions during the application process provide more useful information for predicting future participation than personal background characteristics. Since behavioral factors are more easily influenced through institutional support and communication strategies, they provide practical opportunities for improving student engagement.

Pattern 2: Early Application is Associated with Higher Participation

The predictive model consistently identified application timing as one of the strongest indicators of future participation. Students who applied earlier generally received higher participation probabilities and lower predicted dropout risk.

This pattern suggests that students who demonstrate early interest are often more committed to completing the opportunity. Conversely, delayed applications may indicate lower motivation or greater uncertainty, increasing the likelihood of withdrawal before completion.

Encouraging students to apply earlier through timely notifications, awareness campaigns, and simplified application procedures may therefore contribute to improved participation outcomes.

Pattern 3: Opportunity Categories Show Different Participation Behaviour

The analysis also revealed that student participation varies across different opportunity categories. Engagement activities, Events, Competitions, and Courses generally exhibited higher participation probabilities and lower dropout risk, whereas Internship opportunities showed comparatively higher predicted dropout probabilities.

This pattern indicates that different types of opportunities present different participation challenges. Internship programs often involve longer durations, stricter requirements, and greater time commitments, making sustained participation more difficult for some students.

These findings suggest that participation improvement strategies should be tailored according to the characteristics of each opportunity category rather than applying a single approach across all programs.

Pattern 4: Most Students are Predicted to Participate Successfully

The predicted probability distributions indicate that the majority of students fall within the **Very High Participation** category. This demonstrates that the historical dataset is dominated by successful participation records and that the model predicts positive engagement for most students.

However, a smaller group of students consistently appears within the lower participation categories. Although this group represents a minority of the overall population, it is particularly important because these students are more likely to withdraw from opportunities and therefore require additional support and monitoring.

Pattern 5: Machine Learning Enables Early Risk Identification

A significant advantage of the predictive model is its ability to estimate participation risk before an opportunity is completed. Instead of identifying dropout only after it occurs, the model provides probability estimates that highlight students who may require early intervention.

This proactive approach allows institutions to move from reactive decision-making to preventive support strategies. By identifying at-risk students at an early stage, educators can provide reminders, mentoring, counseling, or additional engagement activities before participation declines.

Overall Summary of Patterns

Collectively, these findings demonstrate that student participation is influenced by a combination of behavioral actions and opportunity characteristics rather than by demographic information alone. The machine learning model successfully captured these relationships and translated them into meaningful probability estimates that support evidence-based decision-making.

These patterns not only validate the effectiveness of the predictive model but also provide valuable guidance for designing future participation strategies aimed at improving student engagement and reducing dropout risk.

The patterns identified through predictive modeling provide a strong foundation for developing practical recommendations. The following section explains how these findings can be translated into actionable strategies that support improved participation, targeted interventions, and more effective opportunity management in Week 4.

4.7 Insights for Week 4 Recommendations

The predictive analysis conducted in Week 3 provides valuable evidence that can be used to develop meaningful recommendations for improving student participation and reducing dropout risk. Rather than relying on assumptions or general observations, the recommendations presented in the next phase of the project will be based directly on the patterns identified through machine learning predictions, feature importance analysis, and probability estimation.

The insights obtained from the Random Forest model highlight areas where educational institutions and opportunity providers can focus their efforts to improve student engagement. By understanding the factors that influence participation and identifying students who are more likely to disengage, organizations can design targeted interventions that are both efficient and data-driven.

Insight 1: Prioritize Early Student Engagement

One of the strongest findings from the predictive model is the importance of student application behavior. Since **Days to Apply** emerged as the most influential feature, encouraging students to apply earlier may contribute to stronger participation and lower dropout risk.

Future recommendations may therefore include sending timely application reminders, increasing awareness of available opportunities, simplifying application procedures, and providing clear guidance during the application process. These initiatives can motivate students to engage with opportunities at an earlier stage, improving their likelihood of successful participation.

Insight 2: Focus Additional Support on High-Risk Opportunity Categories

The predicted probability analysis revealed that not all opportunity categories exhibit the same level of participation strength. Internship opportunities demonstrated comparatively higher predicted dropout probabilities than Engagement activities, Events, Competitions, and Courses.

Based on these findings, future recommendations may focus on strengthening support mechanisms for internship participants. Possible strategies include regular progress monitoring, mentor support, periodic communication, and early identification of students who may require additional guidance throughout the internship period.

Rather than applying identical engagement strategies across all opportunity types, institutions can allocate resources according to the level of predicted participation risk associated with each category.

Insight 3: Use Predictive Analytics for Early Intervention

One of the most valuable outcomes of the machine learning model is its ability to estimate participation probabilities before actual dropout occurs. This enables institutions to move beyond reactive decision-making and adopt a proactive approach to student support.

Students identified with lower participation probabilities can be monitored more closely and offered personalized assistance before disengagement becomes a significant issue. Early interventions such as counseling, reminder notifications, academic guidance, or additional communication may help improve participation outcomes and reduce the likelihood of withdrawal.

Insight 4: Develop Data-Driven Participation Strategies

The feature importance analysis demonstrates that behavioral variables provide stronger predictive information than demographic characteristics. This suggests that future engagement strategies should focus on improving student behavior during the participation process rather than relying primarily on demographic segmentation.

Institutions can use historical participation data together with predictive analytics to continuously evaluate student engagement, monitor participation trends, and refine their support strategies over time. A data-driven approach enables decisions to be supported by objective evidence rather than assumptions, leading to more effective participation management.

Insight 5: Continue Improving the Predictive Model

Although the Random Forest classifier demonstrated strong predictive performance, machine learning models can always be improved as additional data becomes available. Future work may include incorporating new behavioral variables, academic performance indicators, attendance records, or interaction history to enhance prediction accuracy.

Further improvements such as hyperparameter tuning, cross-validation, or experimenting with advanced algorithms like Gradient Boosting or XGBoost may also increase model performance and provide even more reliable participation predictions.

Continuous refinement of the predictive model will enable educational institutions to make increasingly accurate decisions as more historical participation data is collected.

Overall Interpretation

The insights generated from Week 3 demonstrate that predictive analytics has significant potential to support evidence-based educational decision-making. The machine learning model successfully identified the key factors associated with student participation and provided meaningful probability estimates that can guide targeted interventions.

Rather than replacing human decision-making, these predictions should be viewed as decision-support tools that assist educators and administrators in identifying participation trends, prioritizing resources, and implementing timely support for students who may be at greater risk of disengagement.

The recommendations developed in Week 4 will build directly upon these insights, translating the predictive findings into practical strategies aimed at improving student engagement, increasing successful participation, and reducing overall dropout rates.

The predictive insights presented in this report establish a strong analytical foundation for Week 4, where these findings will be transformed into practical recommendations and strategic action plans designed to enhance student participation and support long-term engagement.

Predictive Finding	Recommendation Direction
Early applications reduce dropout risk	Encourage timely applications through reminders and awareness campaigns.
Internship opportunities have higher dropout probability	Strengthen mentoring and monitoring during internships.
Behavioral features are the strongest predictors	Focus interventions on student engagement and application behavior.
Probability scores identify at-risk students	Use predictions to provide early support and personalized guidance.
Random Forest provides reliable predictions	Integrate predictive analytics into regular participation monitoring.

4.8 Conclusion

This week's work focused on applying machine learning techniques to predict student participation outcomes and identify the factors that influence engagement across different opportunity categories. Building upon the cleaned and explored dataset prepared in the previous weeks, predictive models were developed and evaluated to estimate the likelihood of student participation and potential dropout.

Among the models evaluated, the **Random Forest Classifier** demonstrated the strongest predictive performance, outperforming the baseline Logistic Regression model across multiple evaluation metrics. The model achieved higher classification accuracy while effectively capturing the complex relationships between student behavior, opportunity characteristics, and participation outcomes. The inclusion of additional evaluation measures such as Precision, Recall, F1-score, and ROC-AUC provided a comprehensive assessment of the model's performance and confirmed its suitability for predictive analysis.

Feature importance analysis further enhanced the interpretability of the model by identifying the variables that contributed most significantly to prediction accuracy. Behavioral features, particularly

Days to Apply and **Apply to Start Gap**, emerged as the strongest predictors of participation outcomes, indicating that student actions during the application process play a more influential role than demographic characteristics. These findings emphasize the importance of encouraging timely applications and maintaining student engagement throughout the participation journey.

The predicted probability analysis demonstrated that the majority of students are expected to participate successfully, while also identifying a smaller group of students who may be at greater risk of disengagement. The probability estimates generated by the Random Forest model provide valuable opportunities for early intervention by enabling institutions to identify at-risk students before participation declines. In addition, variations observed across opportunity categories highlight the need for category-specific engagement strategies rather than adopting a uniform approach for all opportunities.

The visualizations developed throughout this report—including feature importance charts, confusion matrices, ROC curves, probability distributions, and heatmaps—provided valuable support for interpreting the predictive models and understanding student participation patterns. Together, these visual analyses improved the transparency of the machine learning process and demonstrated how predictive analytics can transform historical participation data into actionable insights.

Overall, the predictive modeling phase successfully achieved the objectives of Week 3 by developing reliable machine learning models, generating meaningful participation probability estimates, evaluating model performance through multiple visualization techniques, and interpreting the resulting findings. The knowledge gained from this analysis establishes a strong analytical foundation for the final phase of the project.

In **Week 4**, these predictive insights will be translated into practical recommendations and strategic action plans aimed at improving student participation, strengthening engagement across opportunity categories, reducing dropout risk, and supporting data-driven decision-making within educational institutions.

Google Collab Link:

https://colab.research.google.com/drive/10YXlh8Sw_JWfCbLRG8SW4CcMutucR_Km?usp=sharing